

AI Governance and Risk Management

From Compliance to Competitive Advantage

Author: Prof. Hernan Huwylar, MBA CPA

Why AI governance must change

Most organizations still treat AI governance as a documentation exercise. Real governance is an operating system that aligns business objectives, risk decisions, engineering practices, and continuous monitoring.

Practical principles

Start with business value, define measurable risks, assign ownership, automate evidence collection, and monitor models after deployment.

Artificial intelligence is changing how organizations make decisions, serve customers, detect fraud, optimize operations, and create products.

Yet many governance programs still resemble traditional compliance initiatives built around static policies and annual reviews.

That approach breaks down once AI systems begin learning from data, interacting with users, or acting autonomously.

Effective AI governance is not about slowing innovation.

It is about creating the confidence to scale it.

Boards increasingly ask the same questions.

Can we trust this model? Who owns the risk? What evidence proves the system remains reliable after deployment? These questions cannot be answered by a policy alone.

They require measurable controls, assigned ownership, continuous monitoring, and operational evidence.

Risk management should create value rather than merely protect assets.

The best programs improve decisions before they reduce losses.

They prioritize opportunities, quantify uncertainty, and help leadership allocate capital where technology creates the greatest return.

AI governance should follow the same philosophy.

#	Date	Article	Why it's worth reading
1	Jul 30, 2026	A Practical Guide for Engineers, Architects, and Governance Teams Who Need to Get It Right	Argues AI security needs more than a repurposed cybersecurity checklist, and shows what the extra layer looks like.
2	Jul 28, 2026	The Seven Gates Every AI Agent Must Clear Before It Can Act (And Most Skip at Least Three)	A control framework for the gap between an agent's decision and its real-world consequence, identity, authority, evidence, execution, verification, obligations, and compensation.
3	Jul 21, 2026	Your Vendor's "We Don't Train On Your Data" Promise Is a Sentence, Not A Data Architecture	Shows where AI vendor data exposure actually lives, fine-tuning, logs, and retrieval, not in the one line procurement teams quote back to legal.
4	Jun 28, 2026	How ISO 24970 and prEN 18229-1 Turn Post-Deployment Chaos Into Auditable Evidence	Connects AI system logging standards to what a board or regulator actually asks for the first time an AI system fails visibly.
5	Jun 27, 2026	How to Build a Policy Engine for AI Agents Without Losing Control	Explains why a system prompt is not a governance control and what a real policy engine for agents needs instead.
6	Jun 17, 2026	The prEN 18286 Reality Check: Ditch Generic AI Governance	A test for whether your AI quality management evidence would survive a notified body or auditor asking to see it directly.
7	May 8, 2026	The prEN 18228 Problem: Why Your AI Risk Assessment Will Fail the First Real Test	Walks through the new European AI risk assessment standard and where most existing assessments fall short of it.
8	Mar 31, 2026	Guide to AI Agent Risk and Control Management Across the Full Lifecycle	End-to-end risk and control coverage for agentic AI, from the first API call an agent makes onward.
9	Mar 30, 2026	Shadow AI Risk Management for CAIOs	A practical approach to the AI tools already inside your organization that nobody

#	Date	Article	Why it's worth reading
			approved, from browser extensions to pasted customer data.
10	Mar 28, 2026	How to Actually Use ISO/IEC 23894 for AI Risk Management	Turns ISO/IEC 23894 into a working process instead of a policy document nobody references again.
11	Mar 16, 2026	What a Chief AI Officer Actually Owns, and What Should Stay With Risk, Legal, and IT	Draws the actual boundary of the CAIO role so it isn't a figurehead title with no real authority.
12	Mar 16, 2026	The AI Use Case Identification and Prioritization Framework	A method for ranking AI ideas by business value before months get spent building the wrong one.
13	Mar 16, 2026	Rules for AI Use, Accountability, BYOAI, Safety by Design, and Content Provenance	Makes the case for six separate AI policies instead of one document trying to serve every audience at once.
14	Mar 16, 2026	Responsible AI Policy Categories	Converts vague principles like "we value fairness" into policy language with an actual operational requirement behind it.
15	Mar 16, 2026	Ways to Calculate Automation Savings and Revenue in AI Projects	Shows how to turn "this will save time" into numbers finance and the executive team will actually trust.
16	Mar 16, 2026	How to Build an AI Roadmap That Delivers Value, Controls Risk, and Survives Change	Turns a pile of competing AI ideas into a sequenced roadmap instead of a budget fight.
17	Mar 15, 2026	How to Negotiate AI Agreements That Protect Data, Value, and Liability	Flags the contract issues generic SaaS terms never anticipated, output ownership, training use, and hallucination liability.
18	Mar 15, 2026	AI Governance From Compliance Tasks to Operations	Argues governance stops being a policy conversation the moment agents start changing tickets and calling tools at machine speed.

#	Date	Article	Why it's worth reading
19	Mar 15, 2026	The AI Career Edge Nobody Talks About	On what separates people who actually break into AI roles beyond the standard degree-and-apply path.
20	Mar 15, 2026	AI Threat and Vulnerability Assessment	A STRIDE-based threat modeling guide covering the attack surface a standard security scan misses entirely.
21	Mar 15, 2026	Field Guide to the 8 Factors That Determine Success or Failure of AI Projects	Distills what actually separates AI projects that deliver value from ones that don't, across strategy, people, and process.
22	Mar 15, 2026	Managing AI Projects With Agile, Exploration, and MLOps	Explains why treating an AI project like standard deterministic software development sets it up to fail.
23	Mar 15, 2026	Data and Tool Infrastructure for AI Projects	Covers the gap between a model that works in a notebook and one that can actually reach production.
24	Mar 15, 2026	The Model Robustness and Monitoring Playbook	Uses a real case of a credit model's accuracy quietly drifting to show why post-deployment monitoring can't be an afterthought.
25	Mar 15, 2026	Modeling Practices for Regulated AI	A validation framework built to satisfy data scientists and regulators at the same time, anchored in CFPB adverse-action guidance.
26	Mar 15, 2026	Effective Fixes for Why Data Science Projects Fail	Traces data science failure back to its earliest, most preventable causes, a vague question and the wrong optimized metric.
27	Mar 15, 2026	AI Deployment Governance for Feedback Loops and MLOps	On why strong models still produce weak outcomes when nobody owns the handoff from user feedback to production change.
28	Mar 13, 2026	Managing AI Development and Deployment Projects	Ten practices that separate AI projects that actually ship from the ones that stall indefinitely.

#	Date	Article	Why it's worth reading
29	Mar 13, 2026	AI Performance Auditing	Points out that most AI audits stop at approval and never check what the system is actually doing in live operation.
30	Mar 12, 2026	How to Explain AI Risk Models So Regulators Actually Trust Them	On building explainability into a model before deployment instead of scrambling for it once an auditor asks.
31	Mar 12, 2026	Predictive Risk Model That Makes the Fewest Expensive Mistakes	Reframes model evaluation around the cost of specific errors, not a single flattering accuracy number.
32	Mar 12, 2026	The Risk and Compliance Automation Playbook	On moving from sampling a handful of transactions to monitoring all of them continuously.
33	Mar 12, 2026	AI Risk Modeling Beyond "Is AI Accurate?"	Argues accuracy answers roughly 15 percent of what can actually go wrong with an AI system, and covers the rest.
34	Mar 12, 2026	Machine Learning for Advanced Predictive Risk Modeling	Makes the case that risk teams who can't build predictive models will eventually be replaced by software that can.
35	Mar 12, 2026	Practical Post-Deployment Maintenance for AI Systems	On why most teams over-invest in reaching deployment and under-invest in what keeps a system safe afterward.
36	Mar 12, 2026	Model Selection and Validation for AI Projects	Covers the overfitting-versus-underfitting tradeoff and how to actually prove a chosen model works.
37	Mar 12, 2026	The Step-by-Step AI Integration Playbook	On why an AI feature can work perfectly in isolation while the business around it still feels fragmented.
38	Mar 12, 2026	AI Contract Clauses That Reduce Vendor, Data, and Liability Risk	Pulls procurement, legal, security, and compliance into one disciplined view of AI purchasing instead of four separate ones.

#	Date	Article	Why it's worth reading
39	Mar 12, 2026	Practical AI Service Level Agreements	On writing SLA terms that account for models drifting and degrading quietly, not just uptime.
40	Mar 12, 2026	AI Model Cards That Improves Transparency, Governance, and Real-World Use	Explains why most model cards are written after the fact for compliance theater, and what a useful one looks like instead.
41	Mar 12, 2026	My New Book: AI Management Systems	Introduces his book on turning AI governance from an aspirational statement into an auditable operational system.
42	Mar 12, 2026	How to Monitor AI Systems After Go-Live Without Creating Audit Theater	On catching the quiet drift, stalled adoption, and hidden cost creep that precede most AI project failures.
43	Mar 12, 2026	Why Separating Your AI Build Team From Your AI Ops Team Guarantees Failure	Applies "you build it, you run it" ownership to AI systems where handoffs otherwise get lost between four different teams.
44	Mar 12, 2026	How to Build the Right AI Delivery Team	Names the seven roles an AI team actually needs and the management functions teams routinely forget to assign.
45	Mar 12, 2026	Resource Estimation for AI Projects	Fifteen cost categories that explain why AI budgets overrun even without one single large surprise expense.
46	Mar 12, 2026	Practical AI Assessments	A nine-stage framework answering three questions every AI project has to face: can we build it, are we ready, did it work.
47	Mar 12, 2026	Goal Setting for AI Projects	On defining scope and success in measurable business terms before different stakeholders quietly assume different goals.
48	Mar 12, 2026	Building vs Buying Decisions for AI Systems	Cuts through the build-versus-buy debate before each department's preferred answer gets baked in as the decision.

#	Date	Article	Why it's worth reading
49	Mar 12, 2026	Feasibility Assessment for AI Projects	On testing data quality, model fit, and integration before a proof of concept gets greenlit on enthusiasm alone.
50	Mar 12, 2026	Problem Definition for AI Projects and Use Cases	Points out that most AI projects go wrong before anyone touches a model, in an underspecified problem statement.
51	Mar 12, 2026	Practical Post-Market Monitoring for AI Systems	Argues most AI governance is strongest before launch and weakest exactly when live users create the real risk.
52	Mar 12, 2026	Compliance Controls for AI Systems	Shows the gap between an AI compliance program that looks complete and one that survives its first real incident.
53	Mar 12, 2026	Practical AI Red Team Implementation Tips for Safer, More Resilient AI Systems	Built around a real case of a language model leaking internal training data, with a playbook for red-teaming before that happens to you.
54	Mar 12, 2026	Practical KPI Tracking for AI Projects	On choosing AI project metrics that actually change delivery decisions instead of numbers everyone quietly agrees to ignore.
55	Mar 12, 2026	Implementation Tips for ISO 42005 AI Impact Assessments	On why ISO 42005 assessments fail in the form itself when teams paste in vendor language and skip foreseeable misuse.
56	Mar 12, 2026	Practical Implementation Tips for an AI Fundamental Rights Taxonomy	Built around a 40-page impact assessment that covered accuracy and bias but never once mentioned fundamental rights.
57	Mar 12, 2026	The 49 AI Quality Characteristics That Define Whether Your System Actually Works	A practitioner's walk-through of ISO/IEC 25059, arguing accuracy alone doesn't make an AI system trustworthy or maintainable.

#	Date	Article	Why it's worth reading
58	Mar 12, 2026	The AI Loss Taxonomy Your Risk Assessments Are Missing	Pushes past vague categories like "reputational damage" toward specific, measurable AI loss types.
59	Mar 12, 2026	The 45 AI Threat Vectors That Your Security Team Probably Isn't Tracking	Built around a security team that found seven threat vectors when dozens more actually applied.
60	Mar 12, 2026	The AI Risk Taxonomy Most Organizations Never Build	Argues a 15-line color-coded heat map for "AI bias" is a liability dressed up as a risk register.
61	Mar 12, 2026	A 12-Step Procedure Merging ISO 27005, ISO 23894, ISO 42001, and FAIR	Combines four separate standards into one assessment procedure that produces an actual usable number.
62	Mar 12, 2026	Quantitative Risk Assessment Using Monte Carlo Simulations and Convolution Methods in R	A hands-on, code-backed case for probabilistic modeling over multiplying ordinal likelihood and impact scores together.
63	Mar 12, 2026	Implementation Tips for Expert Calibration and AI-Augmented Risk Estimation	On correcting the anchoring and overconfidence baked into unstructured expert risk estimates.
64	Mar 12, 2026	Practical Implementation Tips for the AI System Lifecycle RACI Matrix	Traces AI governance failures back to their real root cause, nobody owning the problem at the moment it mattered.
65	Mar 12, 2026	Practical Implementation Tips for AI Project Alignment	On why a technically excellent model that isn't tied to a funded, board-approved outcome is just an expensive experiment.
66	Mar 12, 2026	Practical Implementation Tips for a Fundamental Rights Impact Assessment for High-Risk AI Systems	A direct walkthrough of what EU AI Act Article 27 requires deployers to assess, separate from the provider's conformity work.

#	Date	Article	Why it's worth reading
67	Mar 12, 2026	Practical Implementation Tips for Building and Maintaining an AI Compliance Register	Explains why a single AI system triggering obligations across multiple regulatory domains needs its own register, not a generic GRC tool.
68	Mar 12, 2026	Practical AI Compliance Implementation	An AI legal compliance audit program starting from governance and accountability, since every other finding needs an owner.
69	Mar 12, 2026	Practical ISO 42001 Certification Guidance	On why a generic "responsible AI" policy statement fails every audit, and what an operational one requires instead.
70	Mar 12, 2026	I Implemented ISO 42001 For Global Companies	Four specific ways compliance officers get it wrong when they audit a neural network like a standard IT database.
71	Mar 12, 2026	Spent 5 Years Validating Enterprise AI Models	Makes the case that validation is not a one-time sign-off, since a model starts degrading the moment it's deployed.
72	Mar 12, 2026	Career Topics: The Quantitative Risk Architect	His own path into AI risk management, starting from stochastic calculus at ExxonMobil and Deloitte.
73	Mar 7, 2026	Academic and Institutional Affiliations	Background on his academic role at IE University, IE Law School, and IE Business School.